

## Search II



# Roadmap

## Modeling

Modeling Search Problems

## Algorithms

Tree Search

Dynamic Programming

Uniform Cost Search

Programming and Correctness of UCS

A\*

A\* Relaxations

## Learning

Structured Perceptron



## Key idea: state

A **state** is a summary of all the past actions sufficient to choose future actions **optimally**.

past actions (all cities)	1 3 4 6 5 3
state (current city)	1 3 <del>4</del> 6 5 3

# Review



## Definition: search problem

- $s_{\text{start}}$ : starting state
- $\text{Actions}(s)$ : possible actions
- $\text{Cost}(s, a)$ : action cost
- $\text{Succ}(s, a)$ : successor
- $\text{IsEnd}(s)$ : reached end state?


Objective: find the minimum cost path from  $s_{\text{start}}$  to an  $s$  satisfying  $\text{IsEnd}(s)$ .



## Search: uniform cost search correctness



# High-level strategy



- Explored: states we've found the optimal path to
- Frontier: states we've seen, still figuring out how to get there cheaply
- Unexplored: states we haven't seen

# Uniform cost search (UCS)



**Algorithm: uniform cost search [Dijkstra, 1956]**

Add  $s_{\text{start}}$  to **frontier** (priority queue)

Repeat until frontier is empty:

    Remove  $s$  with smallest priority  $p$  from frontier

    If  $\text{IsEnd}(s)$ : return solution

    Add  $s$  to **explored**

    For each action  $a \in \text{Actions}(s)$ :

        Get successor  $s' \leftarrow \text{Succ}(s, a)$

        If  $s'$  already in explored: continue

        Update **frontier** with  $s'$  and priority  $p + \text{Cost}(s, a)$

[live solution: Uniform Cost Search]


# Analysis of uniform cost search



## Theorem: correctness

When a state  $s$  is popped from the frontier and moved to explored, its priority is  $\text{PastCost}(s)$ , the minimum cost to  $s$ .

Proof:





## DP versus UCS

$N$  total states,  $n$  of which are closer than end state

Algorithm	Cycles?	Action costs	Time/space
DP	no	any	$O(N)$
UCS	yes	$\geq 0$	$O(n \log n)$

Note: UCS potentially explores fewer states, but requires more overhead to maintain the priority queue

Note: assume number of actions per state is constant (independent of  $n$  and  $N$ )



# Summary

- Tree search: memory efficient, suitable for huge state spaces but exponential worst-case running time
- State: summary of past actions sufficient to choose future actions optimally
- Graph search: dynamic programming and uniform cost search construct optimal paths (exponential savings!)
- Next: searching faster with  $A^*$ , learning action costs (if time)



think and share

## Question

Suppose we want to travel from city 1 to city  $n$  (going only forward) and back to city 1 (only going backward). It costs  $c_{ij} \geq 0$  to go from  $i$  to  $j$ . Which of the following algorithms can be used to find the minimum cost path (select all that apply)?

depth-first search

breadth-first search

dynamic programming

uniform cost search




Search:  $A^*$




# A\* algorithm


UCS in action:



A\* in action:



# Can uniform cost search be improved?




Problem: UCS orders states by cost from  $s_{\text{start}}$  to  $s$

Goal: take into account cost from  $s$  to  $s_{\text{end}}$

# Exploring states

UCS: explore states in order of  $\text{PastCost}(s)$



Ideal: explore in order of  $\text{PastCost}(s) + \text{FutureCost}(s)$

A\*: explore in order of  $\text{PastCost}(s) + h(s)$



## Definition: Heuristic function

A heuristic  $h(s)$  is any estimate of  $\text{FutureCost}(s)$ .

# A\* search




**Algorithm: A\* search [Hart/Nilsson/Raphael, 1968]**

Run uniform cost search with **modified edge costs**:

$$\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s)$$

Intuition: add a penalty for how much action  $a$  takes us away from the end state

Example:



$$\text{Cost}'(C, B) = \text{Cost}(C, B) + h(B) - h(C) = 1 + (3 - 2) = 2$$




# An example heuristic

Will any heuristic work?

No.

Counterexample:



Doesn't work because of **negative modified edge costs!**

# Consistent heuristics




## Definition: consistency

A heuristic  $h$  is **consistent** if

- $\text{Cost}'(s, a) = \text{Cost}(s, a) + h(\text{Succ}(s, a)) - h(s) \geq 0$
- $h(s_{\text{end}}) = 0$ .

Condition 1: needed for UCS to work (triangle inequality).



Condition 2:  $\text{FutureCost}(s_{\text{end}}) = 0$  so match it.

# Correctness of $A^*$




## Proposition: correctness

If  $h$  is consistent,  $A^*$  returns the minimum cost path.

# Proof of A\* correctness

- Consider any path  $[s_0, a_1, s_1, \dots, a_L, s_L]$ :



- Key identity:

$$\underbrace{\sum_{i=1}^L \text{Cost}'(s_{i-1}, a_i)}_{\text{modified path cost}} = \underbrace{\sum_{i=1}^L \text{Cost}(s_{i-1}, a_i)}_{\text{original path cost}} + \underbrace{h(s_L) - h(s_0)}_{\text{constant}}$$

- Therefore, A\* (finding the minimum cost path using modified costs) solves the original problem (even though edge costs are all different!)

# Efficiency of A\*



## Theorem: efficiency of A\*

A\* explores all states  $s$  satisfying  
 $\text{PastCost}(s) \leq \text{PastCost}(s_{\text{end}}) - h(s)$

Interpretation: the larger  $h(s)$ , the better


Proof: A\* explores all  $s$  such that

$$\text{PastCost}(s) + h(s)$$

$$\leq$$

$$\text{PastCost}(s_{\text{end}})$$

# Amount explored



- If  $h(s) = 0$ , then  $A^*$  is same as UCS.
- If  $h(s) = \text{FutureCost}(s)$ , then  $A^*$  only explores nodes on a minimum cost path.
- Usually  $h(s)$  is somewhere in between.

# A\* search



**Key idea: distortion**

A\* distorts edge costs to favor end states.



# Admissibility



## Definition: admissibility

A heuristic  $h(s)$  is admissible if  
$$h(s) \leq \text{FutureCost}(s)$$

Intuition: admissible heuristics are optimistic



## Theorem: consistency implies admissibility

If a heuristic  $h(s)$  is **consistent**, then  $h(s)$  is **admissible**.

Proof: use induction on  $\text{FutureCost}(s)$





# Search: $A^*$ relaxations



How do we get good heuristics? Just relax...



# Relaxation

Intuition: ideally, use  $h(s) = \text{FutureCost}(s)$ , but that's as hard as solving the original problem.



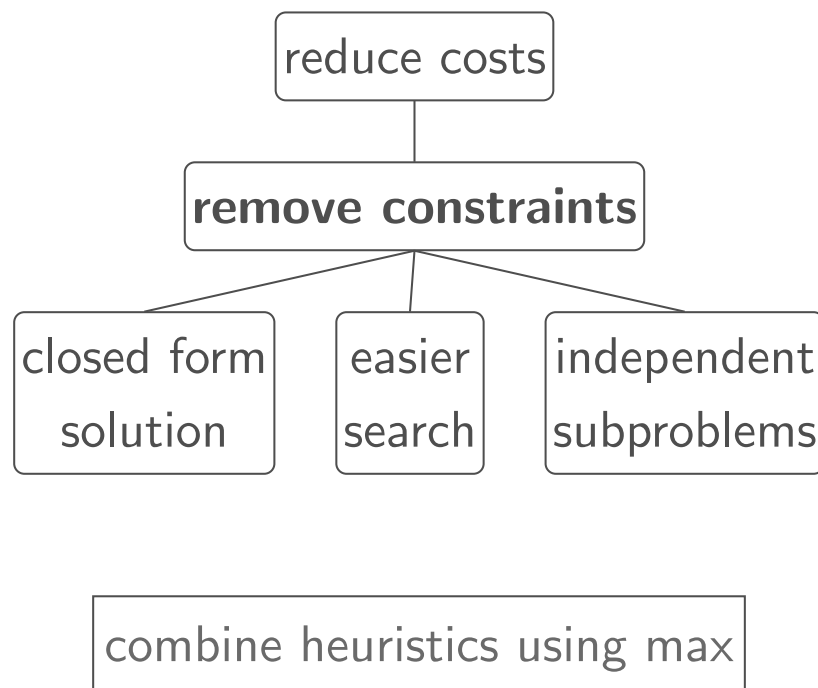
**Key idea: relaxation**

Constraints make life hard. Get rid of them.  
But this is just for the heuristic!





# Relaxation overview




# Closed form solution




Example: knock down walls

Goal: move from triangle to circle



Hard



Easy

Heuristic:

$$h(s) = \text{ManhattanDistance}(s, (2, 5))$$

$$\text{e.g., } h((1, 1)) = 5$$



## Easier search



Example: original problem

Start state: 1

Walk action: from  $s$  to  $s + 1$  (cost: 1)

Tram action: from  $s$  to  $2s$  (cost: 2)

End state:  $n$

**Constraint:** can't have more tram actions than walk actions.

State: (location, **#walk** - **#tram**)

Number of states goes from  $O(n)$  to  $O(n^2)$ !



# Easier search



## Example: relaxed problem

Start state: 1

Walk action: from  $s$  to  $s + 1$  (cost: 1)

Tram action: from  $s$  to  $2s$  (cost: 2)

End state:  $n$

~~Constraint: can't have more tram actions than walk actions.~~

**Original state:** (location, #walk - #tram)

**Relaxed state:** location

# Easier search

- Compute relaxed  $\text{FutureCost}_{\text{rel}}(\text{location})$  for **each** location  $(1, \dots, n)$  using dynamic programming or UCS



## Example: reversed relaxed problem

Start state:  $n$

Walk action: from  $s$  to  $s - 1$  (cost: 1)

Tram action: from  $s$  to  $s/2$  (cost: 2)

End state: 1

Modify UCS to compute all past costs in reversed relaxed problem (equivalent to future costs in relaxed problem!)


- Define heuristic for original problem:

$$h((\text{location}, \# \text{walk} - \# \text{tram})) = \text{FutureCost}_{\text{rel}}(\text{location})$$



# Independent subproblems

[8 puzzle]



Original problem: tiles cannot overlap (constraint)

Relaxed problem: tiles can overlap (no constraint)

Relaxed solution: 8 indep. problems, each in closed form



**Key idea: independence**

Relax original problem into independent subproblems.

# General framework

## Removing constraints


(knock down walls, walk/tram freely, overlap pieces)



## Reducing edge costs

(from  $\infty$  to some finite cost)

Example:



Original:  $\text{Cost}((1, 1), \text{East}) = \infty$

Relaxed:  $\text{Cost}_{\text{rel}}((1, 1), \text{East}) = 1$

# General framework



## Definition: relaxed search problem

A **relaxation**  $P_{\text{rel}}$  of a search problem  $P$  has costs that satisfy:

$$\text{Cost}_{\text{rel}}(s, a) \leq \text{Cost}(s, a).$$



## Definition: relaxed heuristic

Given a relaxed search problem  $P_{\text{rel}}$ , define the **relaxed heuristic**  $h(s) = \text{FutureCost}_{\text{rel}}(s)$ , the minimum cost from  $s$  to an end state using  $\text{Cost}_{\text{rel}}(s, a)$ .

# General framework



## Theorem: consistency of relaxed heuristics

Suppose  $h(s) = \text{FutureCost}_{\text{rel}}(s)$  for some relaxed problem  $P_{\text{rel}}$ .

Then  $h(s)$  is a consistent heuristic.

Proof:

$$h(s) \leq \text{Cost}_{\text{rel}}(s, a) + h(\text{Succ}(s, a)) \text{ [triangle inequality]}$$

$$\leq \text{Cost}(s, a) + h(\text{Succ}(s, a)) \text{ [relaxation]}$$

# Tradeoff

## Efficiency:

$h(s) = \text{FutureCost}_{\text{rel}}(s)$  must be easy to compute

Closed form, easier search, independent subproblems

## Tightness:

heuristic  $h(s)$  should be close to  $\text{FutureCost}(s)$

Don't remove too many constraints

# Max of two heuristics

How do we combine two heuristics?



**Proposition: max heuristic**

Suppose  $h_1(s)$  and  $h_2(s)$  are consistent.


Then  $h(s) = \max\{h_1(s), h_2(s)\}$  is consistent.

Proof: exercise



## Search: recap





# Modeling: Transportation example



## Example: transportation

Street with blocks numbered 1 to  $n$ .

Walking from  $s$  to  $s + 1$  takes 1 minute.

Taking a magic tram from  $s$  to  $2s$  takes 2 minutes.

How to travel from 1 to  $n$  in the least time?



# Inference

## Algorithms

Tree Search

Dynamic Programming

Uniform Cost Search

Programming and Correctness of UCS

A\*

A\* Relaxations

# Dynamic programming



Key idea: state

A **state** is a summary of all the past actions sufficient to choose future actions **optimally**.


past actions (all cities)      1 3 4 6

state (current city)          1 3 4 6


# A\* algorithm

Add in heuristic estimate of future costs.

UCS in action:



A\* in action:



How do we get good heuristics? Just relax...



# Relaxation (breaking the rules)

A framework for producing consistent heuristics.




**Key idea: relaxation**

Constraints make life hard. Get rid of them.  
But this is just for the heuristic!



# Outlook: Sampling Based Planning Algorithms



Probabilistic Roadmaps (PRM) and Rapidly exploring Random Trees (RRT)

# Next time: MDPs




When actions have unknown consequences...



## Search: structured perceptron [optional]







# Search

Transportation example

Start state: 1

Walk action: from  $s$  to  $s + 1$  (cost: 1)

Tram action: from  $s$  to  $2s$  (cost: 2)


End state:  $n$



search algorithm

walk walk tram tram tram walk tram tram

(minimum cost path)



# Learning

Transportation example

Start state: 1

Walk action: from  $s$  to  $s + 1$  (cost: ?)

Tram action: from  $s$  to  $2s$  (cost: ?)

End state:  $n$

walk walk tram tram tram walk tram tram



learning algorithm

walk cost: **1**, tram cost: **2**

# Learning as an inverse problem

Forward problem (search):


$$\text{Cost}(s, a) \longrightarrow (a_1, \dots, a_k)$$

Inverse problem (learning):

$$(a_1, \dots, a_k) \longrightarrow \text{Cost}(s, a)$$

# Prediction (inference) problem

Input  $x$ : search problem without costs




Output  $y$ : solution path

walk walk walk

# Tweaking costs


Costs: {walk:3, tram:2}

Minimum cost path:



Costs: {walk:1, tram:3}

Minimum cost path:




# Modeling costs (simplified)

Assume costs depend only on the action:

$$\text{Cost}(s, a) = \mathbf{w}[a]$$

Candidate output path:



Path cost:

$$\text{Cost}(y) = \mathbf{w}[a_1] + \mathbf{w}[a_2] + \mathbf{w}[a_3]$$

# Learning algorithm



## Algorithm: Structured Perceptron (simplified)

- For each action:  $\mathbf{w}[a] \leftarrow 0$
- For each iteration  $t = 1, \dots, T$ :
  - For each training example  $(x, y) \in \mathcal{D}_{\text{train}}$ :
    - Compute the minimum cost path  $y'$  given  $\mathbf{w}$
    - For each action  $a \in y$ :  $\mathbf{w}[a] \leftarrow \mathbf{w}[a] - 1$
    - For each action  $a \in y'$ :  $\mathbf{w}[a] \leftarrow \mathbf{w}[a] + 1$
- Try to decrease cost of true  $y$  (from training data)
- Try to increase cost of predicted  $y'$  (from search)


[live solution: Structured Perceptron]

# Generalization to features (skip)

Costs are parametrized by feature vector:

$$\text{Cost}(s, a) = \mathbf{w} \cdot \phi(s, a)$$

Example:



Path cost:

$$\text{Cost}(y) = 2 + 1 = 3$$



# Learning algorithm (skip)



## Algorithm: Structured Perceptron [Collins, 2002]

- For each action:  $\mathbf{w} \leftarrow 0$
- For each iteration  $t = 1, \dots, T$ :
  - For each training example  $(x, y) \in \mathcal{D}_{\text{train}}$ :
    - Compute the minimum cost path  $y'$  given  $\mathbf{w}$
    - $\mathbf{w} \leftarrow \mathbf{w} - \phi(y) + \phi(y')$
- Try to decrease cost of true  $y$  (from training data)
- Try to increase cost of predicted  $y'$  (from search)

# Applications

- Part-of-speech tagging

*Fruit flies like a banana.* —————> Noun Noun Verb Det Noun

- Machine translation

*la maison bleue* —————> *the blue house*

# Homework

due: next week

作业 3-周3-torch-nn