


MDPs 2: Reinforcement Learning



Review: MDPs



Definition: Markov decision process

States: the set of states

$s_{\text{start}} \in \text{States}$: starting state

$\text{Actions}(s)$: possible actions from state s

$T(s, a, s')$: probability of s' if take action a in state s

$\text{Reward}(s, a, s')$: reward for the transition (s, a, s')

$\text{IsEnd}(s)$: whether at end of game

$0 \leq \gamma \leq 1$: discount factor (default: 1)

Review: MDPs

- Following a **policy** π produces a path (**episode**)

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

- **Value** function $V_\pi(s)$: expected utility if follow π from state s

$$V_\pi(s) = \begin{cases} 0 & \text{if IsEnd}(s) \\ Q_\pi(s, \pi(s)) & \text{otherwise.} \end{cases}$$

- **Q-value** function $Q_\pi(s, a)$: expected utility if first take action a from state s and then follow π

$$Q_\pi(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_\pi(s')]$$

Unknown transitions and rewards



Definition: Markov decision process

States: the set of states

$s_{\text{start}} \in \text{States}$: starting state

$\text{Actions}(s)$: possible actions from state s

$\text{IsEnd}(s)$: whether at end of game

$0 \leq \gamma \leq 1$: discount factor (default: 1)

reinforcement learning!

CS234 course: <https://web.stanford.edu/class/cs234/>


Mystery game



Example: mystery buttons

For each round $r = 1, 2, \dots$

- You choose A or B.
- You move to a new state and get some rewards.



From MDPs to reinforcement learning



Markov decision process (offline)

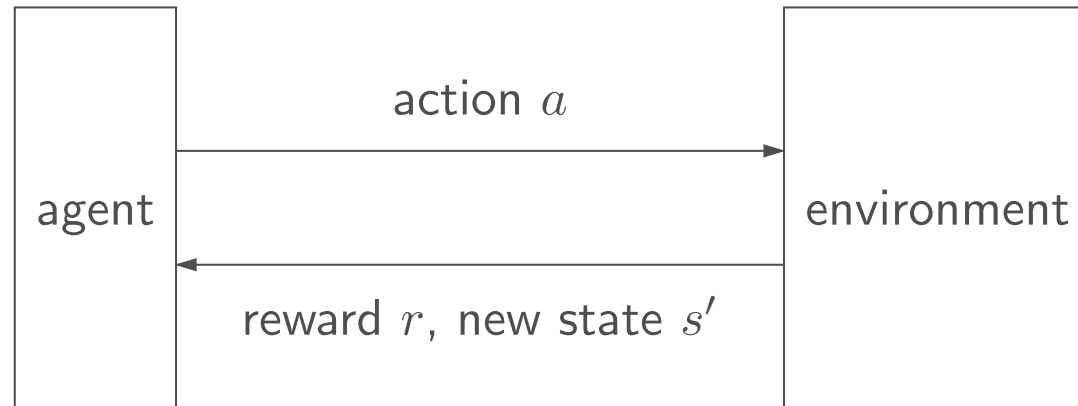
- Have mental model of how the world works.
- Find policy to collect maximum rewards.



Reinforcement learning (online)

- Don't know how the world works.
- Perform actions in the world to find out and collect rewards.

Reinforcement learning framework



Algorithm: reinforcement learning template

For $t = 1, 2, 3, \dots$

Choose action $a_t = \pi_{\text{act}}(s_{t-1})$ (**how?**)

Receive reward r_t and observe new state s_t

Update parameters (**how?**)

Volcano crossing



Run (or press ctrl-enter)

		-50	20
		-50	
2			

Utility: 2

a r s
 (2,1)
 W 0 (2,1)
 W 0 (2,1)
 N 0 (1,1)
 W 0 (1,1)
 N 0 (1,1)
 E 0 (1,2)
 S 0 (2,2)
 W 0 (2,1)
 N 0 (2,2)
 N 0 (3,2)
 S 0 (3,2)
 W 2 (3,1)



Outline

MDPs: model-based methods

MDPs: model-free methods

MDPs: SARSA

MDPs: Q-learning

MDPs: epsilon-greedy

MDPs: function approximation

MDPs: recap and extensions

Model-Based Value Iteration

Data: $s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$



Key idea: model-based learning

Estimate the MDP: $T(s, a, s')$ and $\text{Reward}(s, a, s')$


Transitions:

$$\hat{T}(s, a, s') = \frac{\# \text{ times } (s, a, s') \text{ occurs}}{\# \text{ times } (s, a) \text{ occurs}}$$

Rewards:

$$\widehat{\text{Reward}}(s, a, s') = r \text{ in } (s, a, r, s')$$

Model-Based Value Iteration




Data (following policy $\pi(s) = \text{stay}$):

[in; stay, 4, end]

- Estimates converge to true values (under certain conditions)
- With estimated MDP $(\hat{T}, \widehat{\text{Reward}})$, compute policy using value iteration

Problem



Problem: won't even see (s, a) if $a \neq \pi(s)$ ($a = \text{quit}$)



Key idea: exploration

To do reinforcement learning, need to explore the state space.

Solution: need π to **explore** explicitly (more on this later)



Outline

MDPs: model-based methods

MDPs: model-free methods

MDPs: SARSA

MDPs: Q-learning

MDPs: epsilon-greedy

MDPs: function approximation

MDPs: recap and extensions

From model-based to model-free

$$\hat{Q}_{\text{opt}}(s, a) = \sum_{s'} \hat{T}(s, a, s') [\widehat{\text{Reward}}(s, a, s') + \gamma \hat{V}_{\text{opt}}(s')]$$

All that matters for prediction is (estimate of) $Q_{\text{opt}}(s, a)$.



Key idea: model-free learning

Try to estimate $Q_{\text{opt}}(s, a)$ directly.

Model-free Monte Carlo

Data (following policy π):

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

Recall:

$Q_\pi(s, a)$ is expected utility starting at s , first taking action a , and then following policy π

Utility:


$$u_t = r_t + \gamma \cdot r_{t+1} + \gamma^2 \cdot r_{t+2} + \dots$$

Estimate:

$$\hat{Q}_\pi(s, a) = \text{average of } u_t \text{ where } s_{t-1} = s, a_t = a$$

(and s, a doesn't occur in s_0, \dots, s_{t-2})

Model-free Monte Carlo



Data (following policy $\pi(s) = \text{stay}$):

[in; stay, 4, in; stay, 4, in; stay, 4, in; stay, 4, end]

Note: we are estimating Q_π now, not Q_{opt}



Definition: on-policy versus off-policy

On-policy: estimate the value of data-generating policy

Off-policy: estimate the value of another policy

Model-free Monte Carlo (equivalences)

Data (following policy π):

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

Original formulation

$$\hat{Q}_\pi(s, a) = \text{average of } u_t \text{ where } s_{t-1} = s, a_t = a$$

Equivalent formulation (convex combination)

On each (s, a, u) :

$$\eta = \frac{1}{1 + (\# \text{ updates to } (s, a))}$$
$$\hat{Q}_\pi(s, a) \leftarrow (1 - \eta)\hat{Q}_\pi(s, a) + \eta u$$

Model-free Monte Carlo (equivalences)

Equivalent formulation (convex combination)

On each (s, a, u) :

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\pi}(s, a) + \eta u$$

Equivalent formulation (stochastic gradient)

On each (s, a, u) :

$$\hat{Q}_{\pi}(s, a) \leftarrow \hat{Q}_{\pi}(s, a) - \eta \left[\underbrace{\hat{Q}_{\pi}(s, a)}_{\text{prediction}} - \underbrace{u}_{\text{target}} \right]$$

Implied objective: least squares regression

$$(\hat{Q}_{\pi}(s, a) - u)^2$$

Volcanic model-free Monte Carlo

Run (or press ctrl-enter)

<div>0</div> <div>1 0</div> <div>1</div>	<div>0</div> <div>0 0</div> <div>0</div>	-50	20
<div>1</div> <div>0 1</div> <div>1</div>	<div>0</div> <div>1 0</div> <div>0</div>	-50	<div>0</div> <div>0 0</div> <div>0</div>
2	<div>0</div> <div>0 0</div> <div>0</div>	<div>0</div> <div>0 0</div> <div>0</div>	<div>0</div> <div>0 0</div> <div>0</div>

$a \quad r \quad s$
 $\cdot \quad (2,1)$
 $N \ 0 \ (1,1)$
 $W \ 0 \ (1,1)$
 $S \ 0 \ (2,1)$
 $E \ 0 \ (2,2)$
 $W \ 0 \ (2,1)$
 $S \ 2 \ (3,1)$

Utility: 2



Outline

MDPs: model-based methods

MDPs: model-free methods

MDPs: SARSA

MDPs: Q-learning

MDPs: epsilon-greedy

MDPs: function approximation

MDPs: recap and extensions

Using the utility

Data (following policy $\pi(s) = \text{stay}$):

[in; stay, 4, end]	$u = 4$
[in; stay, 4, in; stay, 4, end]	$u = 8$
[in; stay, 4, in; stay, 4, in; stay, 4, end]	$u = 12$
[in; stay, 4, in; stay, 4, in; stay, 4, in; stay, 4, end]	$u = 16$



Algorithm: model-free Monte Carlo

On each (s, a, u) :

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\pi}(s, a) + \eta \underbrace{u}_{\text{data}}$$

Using the reward + Q-value

Current estimate: $\hat{Q}_\pi(s, \text{stay}) = 11$

Data (following policy $\pi(s) = \text{stay}$):

[in; stay, 4, end]	$4 + 0$
[in; stay, 4, in; stay, 4, end]	$4 + 11$
[in; stay, 4, in; stay, 4, in; stay, 4, end]	$4 + 11$
[in; stay, 4, in; stay, 4, in; stay, 4, in; stay, 4, end]	$4 + 11$



Algorithm: SARSA

On each (s, a, r, s', a') :

$$\hat{Q}_\pi(s, a) \leftarrow (1 - \eta)\hat{Q}_\pi(s, a) + \eta \left[\underbrace{r}_{\text{data}} + \gamma \underbrace{\hat{Q}_\pi(s', a')}_{\text{estimate}} \right]$$

Model-free Monte Carlo versus SARSA



Key idea: bootstrapping

SARSA uses estimate $\hat{Q}_\pi(s, a)$ instead of just raw data u .

u

based on one path

unbiased

large variance

wait until end to update


$r + \hat{Q}_\pi(s', a')$

based on estimate

biased

small variance

can update immediately



Question

Which of the following algorithms allows you to estimate $Q^*(s, a)$ (select all that apply)?

(a) model-based value iteration

(b) model-free Monte Carlo

(c) SARSA



Outline

MDPs: model-based methods

MDPs: model-free methods

MDPs: SARSA

MDPs: Q-learning

MDPs: epsilon-greedy

MDPs: function approximation

MDPs: recap and extensions

Q-learning

Problem: model-free Monte Carlo and SARSA only estimate Q_π , but want Q_{opt} to act optimally

Output	MDP	reinforcement learning
Q_π	policy evaluation	model-free Monte Carlo, SARSA
Q_{opt}	value iteration	Q-learning

Q-learning

Bellman optimality equation:

$$Q_{\text{opt}}(s, a) = \sum_{s'} T(s, a, s') [\text{Reward}(s, a, s') + \gamma V_{\text{opt}}(s')]$$



Algorithm: Q-learning [Watkins/Dayan, 1992]

On each (s, a, r, s') :

$$\hat{Q}_{\text{opt}}(s, a) \leftarrow (1 - \eta) \underbrace{\hat{Q}_{\text{opt}}(s, a)}_{\text{prediction}} + \eta \underbrace{(r + \gamma \hat{V}_{\text{opt}}(s'))}_{\text{target}}$$

Recall: $\hat{V}_{\text{opt}}(s') = \max_{a' \in \text{Actions}(s')} \hat{Q}_{\text{opt}}(s', a')$

SARSA versus Q-learning



Algorithm: SARSA

On each (s, a, r, s', a') :

$$\hat{Q}_{\pi}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\pi}(s, a) + \eta(r + \gamma\hat{Q}_{\pi}(s', a'))$$



Algorithm: Q-learning [Watkins/Dayan, 1992]

On each (s, a, r, s') :

$$\hat{Q}_{\text{opt}}(s, a) \leftarrow (1 - \eta)\hat{Q}_{\text{opt}}(s, a) + \eta(r + \gamma \max_{a' \in \text{Actions}(s')} \hat{Q}_{\text{opt}}(s', a'))]$$

Volcanic SARSA and Q-learning

Run (or press ctrl-enter)

0 0 0	0 0 0	-50	20
0 0 1	0 0 0	-50	0 0 0
2	0 0 0	0 0 0	0 0 0

a, r, s
(2,1)
S2(3,1)

Utility: 2

Off-Policy versus On-Policy



Definition: on-policy versus off-policy

On-policy: evaluate or improve the data-generating policy

Off-policy: evaluate or learn using data from another policy

	on-policy	off-policy
policy evaluation	Monte Carlo SARSA	
policy optimization		Q-learning

Reinforcement Learning Algorithms

Algorithm	Estimating	Based on
Model-Based Monte Carlo	\hat{T}, \hat{R}	$s_0, a_1, r_1, s_1, \dots$
Model-Free Monte Carlo	\hat{Q}_π	u
SARSA	\hat{Q}_π	$r + \hat{Q}_\pi$
Q-Learning	\hat{Q}_{opt}	$r + \hat{Q}_{\text{opt}}$



Outline

MDPs: model-based methods

MDPs: model-free methods

MDPs: SARSA

MDPs: Q-learning

MDPs: epsilon-greedy

MDPs: function approximation

MDPs: recap and extensions



Exploration



Algorithm: reinforcement learning template

For $t = 1, 2, 3, \dots$

Choose action $a_t = \pi_{\text{act}}(s_{t-1})$ (**how?**)

Receive reward r_t and observe new state s_t

Update parameters (**how?**)

$$s_0; a_1, r_1, s_1; a_2, r_2, s_2; a_3, r_3, s_3; \dots; a_n, r_n, s_n$$

Which **exploration policy** π_{act} to use?

No exploration, all exploitation

Attempt 1: Set $\pi_{\text{act}}(s) = \arg \max_{a \in \text{Actions}(s)} \hat{Q}_{\text{opt}}(s, a)$

Run (or press ctrl-enter)

0 0 0	0 0 0.3	-50	100
0 0 2	0 0.1 -25 2	-50	0 0 0
2	0 2 0.5	0 0 0	0 0 0

$a \ r \ s$
 $(2,1)$
 $E \ 0(2,2)$
 $S \ 0(3,2)$
 $W \ 2(3,1)$

Average (lifetime) utility: 1.95

Problem: $\hat{Q}_{\text{opt}}(s, a)$ estimates are inaccurate, **too greedy!**

No exploitation, all exploration

Attempt 2: Set $\pi_{\text{act}}(s) = \text{random from Actions}(s)$

Run (or press ctrl-enter)

98.4 98.4 98.4	98.4 98.4 98.4	-50	100
98.4 98.4 2	98.4 98.4 98.4	-50	99.2 -49.6 96.2 77.5
2	98.4 2 98.4	-50 98.1 98.4 97.8	98.7 98 98.1 97.9

a, r, s
(2,1)
S2(3,1)

Average (lifetime) utility: -19.15

Problem: average utility is low because exploration is **not guided**

Exploration/exploitation tradeoff



Key idea: **balance**

Need to balance **exploration** and **exploitation**.



Examples from life: restaurants, routes, research

Epsilon-greedy



Algorithm: epsilon-greedy policy

$$\pi_{\text{act}}(s) = \begin{cases} \arg \max_{a \in \text{Actions}} \hat{Q}_{\text{opt}}(s, a) & \text{probability } 1 - \epsilon, \\ \text{random from Actions}(s) & \text{probability } \epsilon. \end{cases}$$

Run (or press ctrl-enter)

99.8 100 100	99.6 100 -50	-50	100
100 100 2	100 100 -50	-50	100 -50 100
2	2 100 100	-50 100 100	100 100 100

$a \quad r \quad s$
 $\cdot \quad (2,1)$
 W 0 (2,1)
 N 0 (1,1)
 S 0 (2,1)
 W 0 (2,1)
 W 0 (2,1)
 E 0 (2,2)
 S 0 (3,2)
 E 0 (3,3)
 E 0 (3,4)
 N 0 (2,4)
 N 100 (1,4)

Average (lifetime) utility: 30.71



Outline

MDPs: model-based methods

MDPs: model-free methods

MDPs: SARSA

MDPs: Q-learning


MDPs: epsilon-greedy

MDPs: function approximation

MDPs: recap and extensions

Generalization

Problem: large state spaces, hard to explore



Average (lifetime) utility: 0.44

Q-learning

Stochastic gradient update:

$$\hat{Q}_{\text{opt}}(s, a) \leftarrow \hat{Q}_{\text{opt}}(s, a) - \eta \left[\underbrace{\hat{Q}_{\text{opt}}(s, a)}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{\text{opt}}(s'))}_{\text{target}} \right]$$

This is **rote learning**: every $\hat{Q}_{\text{opt}}(s, a)$ has a different value

Problem: doesn't generalize to unseen states/actions

Function approximation



Key idea: linear regression model

Define **features** $\phi(s, a)$ and **weights** \mathbf{w} :

$$\hat{Q}_{\text{opt}}(s, a; \mathbf{w}) = \mathbf{w} \cdot \phi(s, a)$$



Example: features for volcano crossing

$$\phi_1(s, a) = \mathbf{1}[a = \text{W}] \qquad \phi_7(s, a) = \mathbf{1}[s = (5, *)]$$

$$\phi_2(s, a) = \mathbf{1}[a = \text{E}] \qquad \phi_8(s, a) = \mathbf{1}[s = (*, 6)]$$

...

...

Function approximation



Algorithm: Q-learning with function approximation

On each (s, a, r, s') :

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \underbrace{[\hat{Q}_{\text{opt}}(s, a; \mathbf{w})]}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{\text{opt}}(s'))}_{\text{target}} \phi(s, a)$$

Implied objective function:

$$\left(\underbrace{\hat{Q}_{\text{opt}}(s, a; \mathbf{w})}_{\text{prediction}} - \underbrace{(r + \gamma \hat{V}_{\text{opt}}(s'))}_{\text{target}} \right)^2$$



Outline

MDPs: model-based methods

MDPs: model-free methods

MDPs: SARSA

MDPs: Q-learning

MDPs: epsilon-greedy

MDPs: function approximation

MDPs: recap and extensions



Summary of MDPs


- **Markov decision processes** (MDPs) cope with uncertainty
- Solutions are **policies** rather than paths
- **Policy evaluation** computes policy value (expected utility)
- **Value iteration** computes optimal value (maximum expected utility) and optimal policy
- Main technique: write recurrences \rightarrow algorithm



Summary of Reinforcement Learning

- Online setting: learn and take actions in the real world!
- Monte Carlo: estimate transitions, rewards, Q-values from data
- Bootstrapping: update towards target that depends on estimate rather than just raw data

Covering the unknown



Epsilon-greedy: balance the exploration/exploitation tradeoff

Function approximation: can generalize to unseen states

Challenges in reinforcement learning

Binary classification (sentiment classification, SVMs):

- Stateless, full supervision

Reinforcement learning (flying helicopters, Q-learning):

- Stateful, partial supervision



Key idea: partial supervision

Reward feedback, but not given the solution directly.



Key idea: state

Rewards depend on previous actions \Rightarrow can have delayed rewards.

States and information

	stateless	state
full supervision	supervised learning (binary classification)	supervised imitation learning (structured prediction)
partial supervision	multi-armed bandits	reinforcement learning

Deep reinforcement learning

just use a neural network for $\hat{Q}_{\text{opt}}(s, a)$, π_{opt} , T , etc


Playing Atari [Google DeepMind, 2013]:



- last 4 frames (images) \Rightarrow 3-layer NN \Rightarrow keystroke
- ϵ -greedy, train over 10M frames with 1M replay memory
- Human-level performance on some games (breakout), less good on others (space invaders)

Deep reinforcement learning


- Policy gradient: train a policy $\pi(a \mid s)$ (say, a neural network) to directly maximize expected reward
- Google DeepMind's AlphaGo (2016), AlphaZero (2017)



- Stanford CS224R course:

<https://cs224r.stanford.edu/>

Applications



Robotics Applications: learning dexterous manipulation, control helicopter to do maneuvers in the air




Backgammon: TD-Gammon plays 1-2 million games against itself, human-level performance



Games: DQN solving Atari Games, openAI Five playing Dota.




Managing datacenters; actions: bring up and shut down machine to minimize time/cost



Routing Autonomous Cars: bring down the total latency of vehicles on the road

Markov decision processes: against nature (e.g., Blackjack)



Next time...

Adversarial games: against opponent (e.g., chess)

