
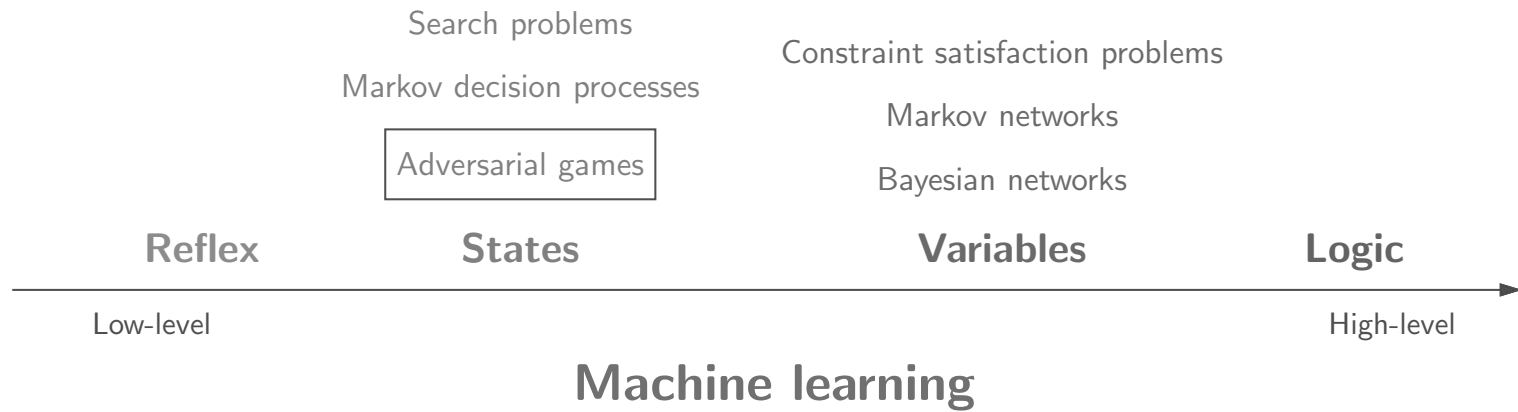


Games I



Course plan



A simple game



Example: game 1

You choose one of the three bins.

I choose a number from that bin.

Your goal is to maximize the chosen number.

A

-50 50

B

1 3

C

-5 15

Roadmap

Modeling

Modeling Games

Algorithms

Game Evaluation

Expectimax

Minimax

Expectiminimax

Evaluation Functions

Alpha-Beta Pruning

Learning

Temporal Difference Learning


Other Topics

Simultaneous Games

Non-Zero-Sum Games



Games: modeling




Game tree



Key idea: game tree

Each node is a decision point for a player.

Each root-to-leaf path is a possible outcome of the game.



Two-player zero-sum games

Players = {agent, opp}



Definition: two-player zero-sum game

s_{start} : starting state

Actions(s): possible actions from state s


Succ(s, a): resulting state if choose action a in state s

IsEnd(s): whether s is an end state (game over)

Utility(s): agent's utility for end state s

Player(s) \in Players: player who controls state s

Example: chess



Players = {white, black}

State s : (position of all pieces, whose turn it is)

Actions(s): legal chess moves that Player(s) can make

IsEnd(s): whether s is checkmate or draw

Utility(s): $+\infty$ if white wins, 0 if draw, $-\infty$ if black wins

Characteristics of games

- All the utility is at the end state



- Different players in control at different states



The halving game



Problem: halving game

Start with a number N .

Players take turns either decrementing N or replacing it with $\lfloor \frac{N}{2} \rfloor$.

The player that is left with 0 wins.

[live solution: HalvingGame]

Policies

Deterministic policies: $\pi_p(s) \in \text{Actions}(s)$

action that player p takes in state s


Stochastic policies $\pi_p(s, a) \in [0, 1]$:

probability of player p taking action a in state s

[live solution: policies, main loop]



Games: game evaluation




Game evaluation example



Example: game evaluation

$$\pi_{\text{agent}}(s) = A$$


$$\pi_{\text{opp}}(s, a) = \frac{1}{2} \text{ for } a \in \text{Actions}(s)$$



$$V_{\text{eval}}(s_{\text{start}}) = 0$$

Game evaluation recurrence

Analogy: recurrence for policy evaluation in MDPs




Value of the game:

$$V_{\text{eval}}(s) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \sum_{a \in \text{Actions}(s)} \pi_{\text{agent}}(s, a) V_{\text{eval}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{agent} \\ \sum_{a \in \text{Actions}(s)} \pi_{\text{opp}}(s, a) V_{\text{eval}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{opp} \end{cases}$$



Games: expectimax




Expectimax example



Example: expectimax


$$\pi_{\text{opp}}(s, a) = \frac{1}{2} \text{ for } a \in \text{Actions}(s)$$



$$V_{\text{exptmax}}(s_{\text{start}}) = 5$$

Expectimax recurrence


Analogy: recurrence for value iteration in MDPs



$$V_{\text{exptmax}}(s) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} V_{\text{exptmax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{agent} \\ \sum_{a \in \text{Actions}(s)} \pi_{\text{opp}}(s, a) V_{\text{exptmax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{opp} \end{cases}$$



Games: minimax



Problem: don't know opponent's policy

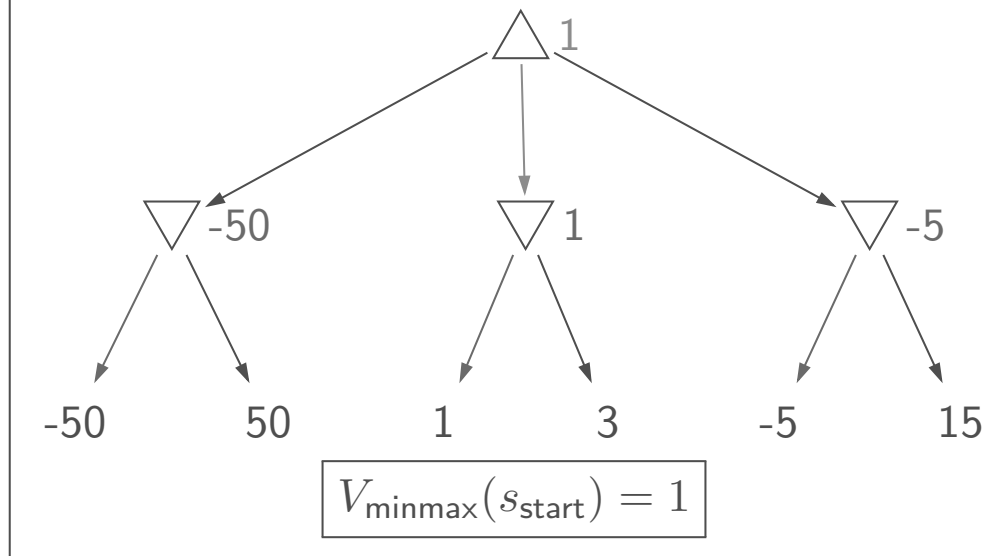
Approach: assume the worst case



Minimax example




Example: minimax



Minimax recurrence

No analogy in MDPs:




$$V_{\text{minmax}}(s) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} V_{\text{minmax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{agent} \\ \min_{a \in \text{Actions}(s)} V_{\text{minmax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{opp} \end{cases}$$

Extracting minimax policies

$$\pi_{\max}(s) = \arg \max_{a \in \text{Actions}(s)} V_{\min\max}(\text{Succ}(s, a))$$

$$\pi_{\min}(s) = \arg \min_{a \in \text{Actions}(s)} V_{\min\max}(\text{Succ}(s, a))$$



The halving game



Problem: halving game

Start with a number N .

Players take turns either decrementing N or replacing it with $\lfloor \frac{N}{2} \rfloor$.

The player that is left with 0 wins.

[live solution: minimaxPolicy]

Face off

Recurrences produce policies:

$$V_{\text{exptmax}} \Rightarrow \pi_{\text{exptmax}(7)}, \pi_7 \text{ (some opponent)}$$

$$V_{\text{minmax}} \Rightarrow \pi_{\text{max}}, \pi_{\text{min}}$$

Play policies against each other:

	π_{min}	π_7
π_{max}	$V(\pi_{\text{max}}, \pi_{\text{min}})$	$V(\pi_{\text{max}}, \pi_7)$
$\pi_{\text{exptmax}(7)}$	$V(\pi_{\text{exptmax}(7)}, \pi_{\text{min}})$	$V(\pi_{\text{exptmax}(7)}, \pi_7)$


What's the relationship between these values?

Minimax property 1



Proposition: best against minimax opponent

$$V(\pi_{\max}, \pi_{\min}) \geq V(\pi_{\text{agent}}, \pi_{\min}) \text{ for all } \pi_{\text{agent}}$$




Minimax property 2



Proposition: lower bound against any opponent

$$V(\pi_{\max}, \pi_{\min}) \leq V(\pi_{\max}, \pi_{\text{opp}}) \text{ for all } \pi_{\text{opp}}$$




Minimax property 3




Proposition: not optimal if opponent is known

$$V(\pi_{\max}, \pi_7) \leq V(\pi_{\text{exptmax}(7)}, \pi_7) \text{ for opponent } \pi_7$$




Relationship between game values



	π_{\min}		π_7
	$V(\pi_{\max}, \pi_{\min})$	\leq	$V(\pi_{\max}, \pi_7)$
π_{\max}	1		2
	\vee		\wedge
	$V(\pi_{\text{exptmax}(7)}, \pi_{\min})$		$V(\pi_{\text{exptmax}(7)}, \pi_7)$
$\pi_{\text{exptmax}(7)}$	-5		5



Games: expectiminimax



A modified game






Example: game 2

You choose one of the three bins.

Flip a coin; if heads, then move one bin to the left (with wrap around).

I choose a number from that bin.

Your goal is to maximize the chosen number.




Expectiminimax example



Example: expectiminimax


$$\pi_{\text{coin}}(s, a) = \frac{1}{2} \text{ for } a \in \{0, 1\}$$



$$V_{\text{exptminmax}}(s_{\text{start}}) = -2$$

Expectiminimax recurrence

Players = {agent, opp, coin}



$$V_{\text{exptminimax}}(s) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \max_{a \in \text{Actions}(s)} V_{\text{exptminimax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{agent} \\ \min_{a \in \text{Actions}(s)} V_{\text{exptminimax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{opp} \\ \sum_{a \in \text{Actions}(s)} \pi_{\text{coin}}(s, a) V_{\text{exptminimax}}(\text{Succ}(s, a)) & \text{Player}(s) = \text{coin} \end{cases}$$



Summary so far

Primitives: max nodes, chance nodes, min nodes

Composition: alternate nodes according to model of game


Value function $V...(\mathbf{s})$: recurrence for expected utility

Scenarios to think about:


- What if you are playing against multiple opponents?
- What if you and your partner have to take turns (table tennis)?
- Some actions allow you to take an extra turn?



Games: evaluation functions



Computation



Approach: tree search

Complexity:

- branching factor b , depth d ($2d$ plies)
- $O(d)$ space, $O(b^{2d})$ time

Chess: $b \approx 35$, $d \approx 50$


25515520672986852924121150151425587630190414488161019324176778440771467258239937365843732987043555789782336195637736653285543297897675074636936187744140625

Speeding up minimax

- Evaluation functions: use domain-specific knowledge, compute approximate answer
- Alpha-beta pruning: general-purpose, compute exact answer



Depth-limited search




Limited depth tree search (stop at maximum depth d_{\max}):

$$V_{\min\max}(s, d) = \begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \text{Eval}(s) & d = 0 \\ \max_{a \in \text{Actions}(s)} V_{\min\max}(\text{Succ}(s, a), d) & \text{Player}(s) = \text{agent} \\ \min_{a \in \text{Actions}(s)} V_{\min\max}(\text{Succ}(s, a), d - 1) & \text{Player}(s) = \text{opp} \end{cases}$$

Use: at state s , call $V_{\min\max}(s, d_{\max})$

Convention: decrement depth at last player's turn

Evaluation functions




Definition: Evaluation function

An evaluation function $\text{Eval}(s)$ is a (possibly very weak) estimate of the value $V_{\min\max}(s)$.

Analogy: $\text{FutureCost}(s)$ in search problems

Evaluation functions



Example: chess

$\text{Eval}(s) = \text{material} + \text{mobility} + \text{king-safety} + \text{center-control}$

$\text{material} = 10^{100}(K - K') + 9(Q - Q') + 5(R - R') +$
 $3(B - B' + N - N') + 1(P - P')$


$\text{mobility} = 0.1(\text{num-legal-moves} - \text{num-legal-moves}')$

...



Summary: evaluation functions


Depth-limited exhaustive search: $O(b^{2d})$ time



- $\text{Eval}(s)$ attempts to estimate $V_{\min\max}(s)$ using domain knowledge
- No guarantees (unlike A^*) on the error from approximation



Games: alpha-beta pruning



Pruning principle

Choose A or B with maximum value:

A: [3, **5**]

B: [**5**, 100]




Key idea: branch and bound

Maintain lower and upper bounds on values.

If intervals don't overlap non-trivially, then can choose optimally without further work.

Pruning game trees



Once we see 2, we know that value of right node must be ≤ 2

Root computes $\max(3, \leq 2) = 3$


Since branch doesn't affect root value, can safely prune

Alpha-beta pruning




Key idea: optimal path

The optimal path is path that minimax policies take.
Values of all nodes on path are the same.



- a_s : lower bound on value of max node s
- b_s : upper bound on value of min node s
- Prune a node if its interval doesn't have non-trivial overlap with every ancestor (store $\alpha_s = \max_{s' \preceq s} a_{s'}$ and $\beta_s = \min_{s' \preceq s} b_{s'}$)


Alpha-beta pruning example



Move ordering

Pruning depends on order of actions.

Can't prune the 5 node:



Move ordering

Which ordering to choose?


- Worst ordering: $O(b^{2 \cdot d})$ time
- Best ordering: $O((\sqrt{b - \frac{3}{4}} + \frac{1}{2})^{2 \cdot d}) \simeq O(b^{2 \cdot 0.5d})$ time
- Random ordering: $O(b^{2 \cdot 0.75d})$ time when $b = 2$
- Random ordering: $O((\frac{b-1+\sqrt{b^2+14b+1}}{4})^{2 \cdot d})$ for general b

In practice, can use evaluation function $\text{Eval}(s)$:

- Max nodes: order successors by decreasing $\text{Eval}(s)$
- Min nodes: order successors by increasing $\text{Eval}(s)$




Games: recap





Summary



- Game trees: model opponents, randomness
- Minimax: find optimal policy against an adversary
- Evaluation functions: domain-specific, approximate
- Alpha-beta pruning: domain-general, exact