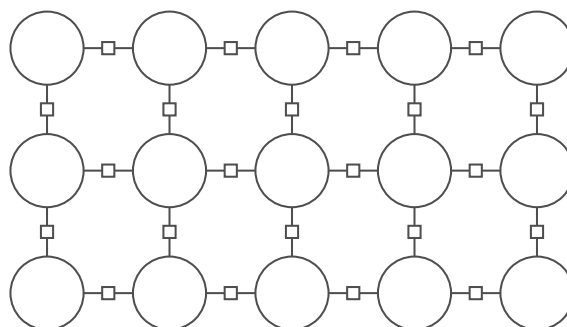


Bayesian Networks II





Lecture: Bayesian networks


Definitions: Probabilistic Programming

Inference: Probabilistic Inference

Inference: Forward Backward

Inference: Particle Filtering

Probabilistic programs



Joint distribution:

$$\mathbb{P}(B = b, E = e, A = a) = p(b)p(e)p(a \mid b, e)$$



Probabilistic program: alarm

$B \sim \text{Bernoulli}(\epsilon)$

$E \sim \text{Bernoulli}(\epsilon)$

$A = B \vee E$

```
def Bernoulli(epsilon):  
    return random.random() < epsilon
```



Key idea: probabilistic program

A randomized program that sets the random variables.

Probabilistic program: example



Probabilistic program: object tracking

$$X_0 = (0, 0)$$


For each time step $i = 1, \dots, n$:

if $\text{Bernoulli}(\alpha)$:

$$X_i = X_{i-1} + (1, 0) \text{ [go right]}$$


else:

$$X_i = X_{i-1} + (0, 1) \text{ [go down]}$$




(press ctrl-enter to save)

Run



Probabilistic inference: example

Question: what are possible trajectories given **evidence** $X_{10} = (8, 2)$?



(press ctrl-enter to save)

Run

Application: language modeling


Can be used to score sentences for speech recognition or machine translation



Probabilistic program: Markov model

For each position $i = 1, 2, \dots, n$:

Generate word $X_i \sim p(X_i \mid X_{i-1})$



Application: object tracking




Probabilistic program: hidden Markov model (HMM)

For each time step $t = 1, \dots, T$:

Generate object location $H_t \sim p(H_t \mid H_{t-1})$

Generate sensor reading $E_t \sim p(E_t \mid H_t)$



Inference: given sensor readings, where is the object?

Application: multiple object tracking




Probabilistic program: factorial HMM

For each time step $t = 1, \dots, T$:

For each object $o \in \{a, b\}$:

Generate location $H_t^o \sim p(H_t^o \mid H_{t-1}^o)$

Generate sensor reading $E_t \sim p(E_t \mid H_t^a, H_t^b)$



Application: document classification




Probabilistic program: naive Bayes

Generate label $Y \sim p(Y)$

For each position $i = 1, \dots, L$:

Generate word $W_i \sim p(W_i \mid Y)$



Inference: given a text document, what is it about?

Application: topic modeling




Probabilistic program: latent Dirichlet allocation

Generate a distribution over topics $\alpha \in \mathbb{R}^K$

For each position $i = 1, \dots, L$:

Generate a topic $Z_i \sim p(Z_i \mid \alpha)$

Generate a word $W_i \sim p(W_i \mid Z_i)$



Inference: given a text document, what topics is it about?

Application: medical diagnosis




Probabilistic program: diseases and symptoms

For each disease $i = 1, \dots, m$:

Generate activity of disease $D_i \sim p(D_i)$

For each symptom $j = 1, \dots, n$:

Generate activity of symptom $S_j \sim p(S_j \mid D_{1:m})$



Inference: If a patient has some symptoms, what diseases do they have?

Application: social network analysis




Probabilistic program: stochastic block model

For each person $i = 1, \dots, n$:

Generate person type $H_i \sim p(H_i)$

For each pair of people $i \neq j$:


Generate connectedness $E_{ij} \sim p(E_{ij} \mid H_i, H_j)$



Inference: Given a social network graph, what types of people are there?



Summary



- Probabilistic program specifies a Bayesian network
- Many different types of models
- Common paradigm: come up with stories of how the quantities of interest (output) generate the data (input)
- Opposite of how we normally do classification!



Lecture: Bayesian networks


Definitions: Probabilistic Programming

Inference: Probabilistic Inference

Inference: Forward Backward

Inference: Particle Filtering

Review: Bayesian network



Random variables:

cold C , allergies A , cough H , itchy eyes I

Joint distribution:

$$\mathbb{P}(C = c, A = a, H = h, I = i) = p(c)p(a)p(h \mid c, a)p(i \mid a)$$




Definition: Bayesian network

Let $X = (X_1, \dots, X_n)$ be random variables.

A **Bayesian network** is a directed acyclic graph (DAG) that specifies a joint distribution over X as a product of local conditional distributions, one for each node:

$$\mathbb{P}(X_1 = x_1, \dots, X_n = x_n) \stackrel{\text{def}}{=} \prod_{i=1}^n p(x_i \mid x_{\text{Parents}(i)})$$

Review: probabilistic inference



Question: $\mathbb{P}(C \mid H = 1, I = 1)$

Input

Bayesian network: $\mathbb{P}(X_1, \dots, X_n)$

Evidence: $E = e$ where $E \subseteq X$ is subset of variables



Query: $Q \subseteq X$ is subset of variables



Output

$\mathbb{P}(Q \mid E = e) \longleftrightarrow \mathbb{P}(Q = q \mid E = e)$ for all values q

Reduction to Markov networks





$$\mathbb{P}(C = c, A = a, H = h, I = i) = \frac{1}{Z} p(c) p(a) p(h \mid c, a) p(i \mid a)$$

Bayesian network = Markov network with normalization constant $Z = 1$

Reminder: single factor that connects **all** parents!

Conditioning on evidence



Markov network:


$$\mathbb{P}(C = c, A = a \mid H = 1, I = 1) = \frac{1}{Z} p(c) p(a) p(h = 1 \mid c, a) p(i = 1 \mid a)$$

Bayesian network with evidence = Markov network with $Z = \mathbb{P}(H = 1, I = 1)$

Solution: run any inference algorithm for Markov networks (e.g., Gibbs sampling)!

[demo]

Leveraging additional structure: unobserved leaves



Markov network:


$$\mathbb{P}(C = c, A = a, I = i \mid H = 1) = \frac{1}{Z} p(c) p(a) p(h = 1 \mid c, a) p(i \mid a),$$

where $Z = \mathbb{P}(H = 1)$

Question: $\mathbb{P}(C = 1 \mid H = 1)$

Can we reduce the Markov network before running inference?

Leveraging additional structure: unobserved leaves




Markov network:

$$\begin{aligned}\mathbb{P}(C = c, A = a \mid H = 1) &= \sum_i \mathbb{P}(C = c, A = a, I = i \mid H = 1) \\ &= \sum_i \frac{1}{Z} p(c) p(a) p(h = 1 \mid c, a) p(i \mid a) \\ &= \frac{1}{Z} p(c) p(a) p(h = 1 \mid c, a) \sum_i p(i \mid a) \\ &= \frac{1}{Z} p(c) p(a) p(h = 1 \mid c, a)\end{aligned}$$

Throw away any unobserved leaves before running inference!

Leveraging additional structure: independence




Markov network:

$$\begin{aligned}\mathbb{P}(C = c \mid I = 1) &= \sum_{a,h} \mathbb{P}(C = c, A = a, H = h \mid I = 1) \\ &= \sum_{a,h} \frac{1}{Z} p(c) p(a) p(h \mid c, a) p(i = 1 \mid a) \\ &= \sum_a \frac{1}{Z} p(c) p(a) p(i = 1 \mid a) \\ &= p(c) \sum_a \frac{1}{Z} p(a) p(i = 1 \mid a) \\ &= p(c)\end{aligned}$$

Throw away any disconnected components before running inference!



Summary



- Condition on evidence (e.g., $I = 1$)
- Throw away unobserved leaves, e.g., I for $\mathbb{P}(C = 1 \mid H = 1)$
- Discard disconnected components, e.g., A and I for $\mathbb{P}(C = c \mid I = 1)$
- Define Markov network out of remaining factors
- Run your favorite inference algorithm (e.g., manual, Gibbs sampling)



Lecture: Bayesian networks



Definitions: Probabilistic Programming

Inference: Probabilistic Inference

Inference: Forward Backward

Inference: Particle Filtering

Hidden Markov models for object tracking




start

H_1

- 1/3
- 1/3
- 1/3


transition

H_i



emission

E_i




h_1	$p(h_1)$
0	1/3
1	1/3
2	1/3

h_i	$p(h_i h_{i-1})$
$h_{i-1} - 1$	1/4
h_{i-1}	1/2
$h_{i-1} + 1$	1/4

e_i	$p(e_i h_i)$
$h_i - 1$	1/4
h_i	1/2
$h_i + 1$	1/4

$$\mathbb{P}(H = h, E = e) = \underbrace{p(h_1)}_{\text{start}} \prod_{i=2}^n \underbrace{p(h_i | h_{i-1})}_{\text{transition}} \prod_{i=1}^n \underbrace{p(e_i | h_i)}_{\text{emission}}$$

Inference questions



Question (**filtering**):

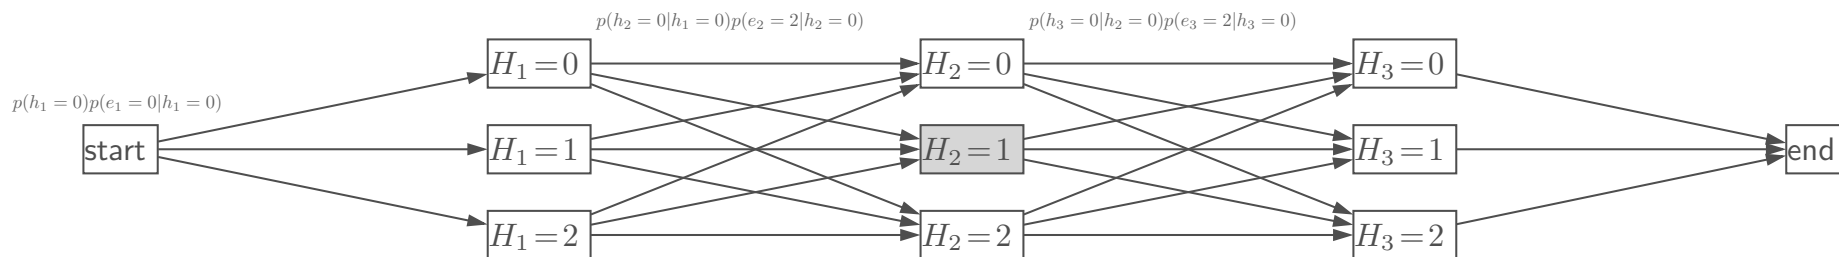
$$\mathbb{P}(H_2 \mid E_1 = 0, E_2 = 2)$$

Question (**smoothing**):

$$\mathbb{P}(H_2 \mid E_1 = 0, E_2 = 2, E_3 = 2)$$

Note: filtering is a special case of smoothing if marginalize unobserved leaves


Lattice representation



- Edge $\boxed{\text{start}} \Rightarrow \boxed{H_1 = h_1}$ has weight $p(h_1)p(e_1 \mid h_1)$
- Edge $\boxed{H_{i-1} = h_{i-1}} \Rightarrow \boxed{H_i = h_i}$ has weight $p(h_i \mid h_{i-1})p(e_i \mid h_i)$
- Each path from $\boxed{\text{start}}$ to $\boxed{\text{end}}$ is an assignment with weight equal to the product of edge weights

Key: $\mathbb{P}(H_i = h_i \mid E = e)$ is the weighted fraction of paths through $\boxed{H_i = h_i}$

Forward and backward messages



Forward: $F_i(h_i) = \sum_{h_{i-1}} F_{i-1}(h_{i-1}) \text{Weight}(\boxed{H_{i-1} = h_{i-1}}, \boxed{H_i = h_i})$

sum of weights of paths from $\boxed{\text{start}}$ to $\boxed{H_i = h_i}$

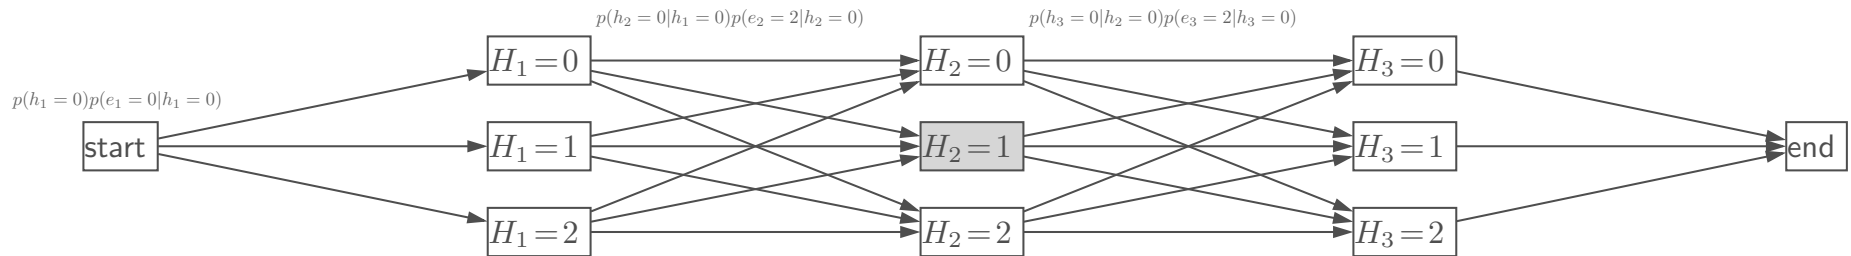
Backward: $B_i(h_i) = \sum_{h_{i+1}} B_{i+1}(h_{i+1}) \text{Weight}(\boxed{H_i = h_i}, \boxed{H_{i+1} = h_{i+1}})$

sum of weights of paths from $\boxed{H_i = h_i}$ to $\boxed{\text{end}}$

Define $S_i(h_i) = F_i(h_i) B_i(h_i)$:

sum of weights of paths from $\boxed{\text{start}}$ to $\boxed{\text{end}}$ through $\boxed{H_i = h_i}$

Putting everything together



$$\mathbb{P}(H_i = h_i \mid E = e) = \frac{S_i(h_i)}{\sum_v S_i(v)}$$



Algorithm: forward-backward algorithm

Compute F_1, F_2, \dots, F_n

Compute B_n, B_{n-1}, \dots, B_1

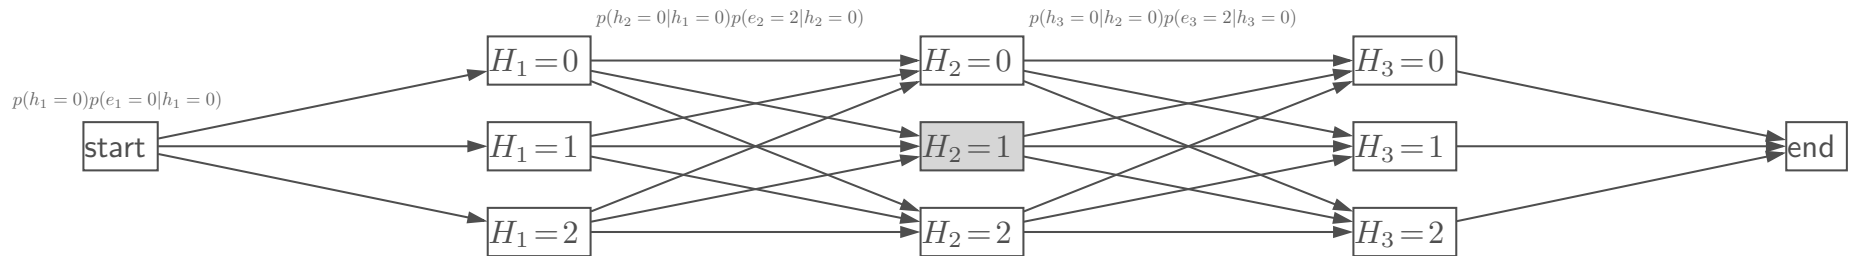
Compute S_i for each i and normalize

Running time: $O(n|\text{Domain}|^2)$

[demo]



Summary



- Lattice representation: paths are assignments
- Dynamic programming: compute sums efficiently
- Forward-backward algorithm: compute all smoothing questions, share intermediate computations



Lecture: Bayesian networks



Definitions: Probabilistic Programming

Inference: Probabilistic Inference

Inference: Forward Backward

Inference: Particle Filtering

Review: Hidden Markov models for object tracking



start


H_1

- 1/3
- 1/3
- 1/3

h_1	$p(h_1)$
0	1/3
1	1/3
2	1/3

transition


H_i



h_i	$p(h_i h_{i-1})$
$h_{i-1} - 1$	1/4
h_{i-1}	1/2
$h_{i-1} + 1$	1/4

emission


E_i



e_i	$p(e_i h_i)$
$h_i - 1$	1/4
h_i	1/2
$h_i + 1$	1/4

$$\mathbb{P}(H = h, E = e) = \underbrace{p(h_1)}_{\text{start}} \prod_{i=2}^n \underbrace{p(h_i | h_{i-1})}_{\text{transition}} \prod_{i=1}^n \underbrace{p(e_i | h_i)}_{\text{emission}}$$

Review: inference in Hidden Markov models




Filtering questions:

$$\mathbb{P}(H_1 \mid E_1 = 0)$$

$$\mathbb{P}(H_2 \mid E_1 = 0, E_2 = 2)$$

$$\mathbb{P}(H_3 \mid E_1 = 0, E_2 = 2, E_3 = 2)$$


Problem: many possible location values for H_i



Forward-backward is too slow ($O(n|\text{Domain}|^2)$)...

Beam search for HMMs

Idea: keep $\leq K$ partial assignments (**particles**)



Algorithm: beam search

Initialize $C \leftarrow [\{\}]$

For each $i = 1, \dots, n$:

Extend:

$$C' \leftarrow \{h \cup \{H_i : v\} : h \in C, v \in \text{Domain}_i\}$$

Prune:


$C \leftarrow K$ particles of C' with highest weights

Normalize weights to get approximate $\hat{\mathbb{P}}(H_1, \dots, H_n \mid E = e)$

Sum probabilities to get any approximate $\hat{\mathbb{P}}(H_i \mid E = e)$

[demo: beamSearch({K:3})]

Beam search problems



Algorithm: beam search

Initialize $C \leftarrow [\{\}]$

For each $i = 1, \dots, n$:

Extend:

$$C' \leftarrow \{h \cup \{H_i : v\} : h \in C, v \in \text{Domain}_i\}$$

Prune:

$$C \leftarrow K \text{ particles of } C' \text{ with highest weights}$$

- Extend: slow because requires considering every possible value for H_i
- Prune: greedily taking best K doesn't provide diversity

Particle filtering solution (3 steps): **propose, weight, resample**

Step 1: propose

Old particles: $\approx \mathbb{P}(H_1, H_2 \mid E_1 = 0, E_2 = 2)$

$\{H_1 : 0, H_2 : 1\}$

$\{H_1 : 1, H_2 : 2\}$



Key idea: proposal distribution

For each old particle (h_1, h_2) , sample $H_3 \sim p(h_3 \mid h_2)$.

h_i	$p(h_i \mid h_{i-1})$
$h_{i-1} - 1$	1/4
h_{i-1}	1/2
$h_{i-1} + 1$	1/4

New particles: $\approx \mathbb{P}(H_1, H_2, H_3 \mid E_1 = 0, E_2 = 2)$

$\{H_1 : 0, H_2 : 1, H_3 : 1\}$

$\{H_1 : 1, H_2 : 2, H_3 : 2\}$

Step 2: weight

Old particles: $\approx \mathbb{P}(H_1, H_2, H_3 \mid E_1 = 0, E_2 = 1)$

$\{H_1 : 0, H_2 : 1 : H_3 : 1\}$

$\{H_1 : 1, H_2 : 2 : H_3 : 2\}$



Key idea: weighting based on evidence

For each old particle (h_1, h_2, h_3) , weight it by $p(e_3 = 2 \mid h_3)$.

h_3	$p(e_3 = 2 \mid h_3)$
0	0
1	1/4
2	1/2

New particles: $\approx \mathbb{P}(H_1, H_2, H_3 \mid E_1 = 0, E_2 = 1, E_3 = 2)$

$\{H_1 : 0, H_2 : 1 : H_3 : 1\} (1/4)$

$\{H_1 : 1, H_2 : 2 : H_3 : 2\} (1/2)$

Step 3: resample

Old particles: $\approx \mathbb{P}(H_1, H_2, H_3 \mid E_1 = 0, E_2 = 2, E_3 = 2)$

$\{H_1 : 0, H_2 : 1 : H_3 : 1\} (1/4) \Rightarrow 1/3$

$\{H_1 : 1, H_2 : 2 : H_3 : 2\} (1/2) \Rightarrow 2/3$



Key idea: resampling

Normalize weights and draw K samples to redistribute particles to more promising areas.


New particles: $\approx \mathbb{P}(H_1, H_2, H_3 \mid E_1 = 0, E_2 = 2, E_3 = 2)$

$\{H_1 : 1, H_2 : 2 : H_3 : 2\}$


$\{H_1 : 1, H_2 : 2 : H_3 : 2\}$

Why sampling?

distribution




K with highest weight



not representative

K sampled from distribution



more representative

Sampling is especially important when there is high uncertainty!

Particle filtering



Algorithm: particle filtering

Initialize $C \leftarrow [\{\}]$

For each $i = 1, \dots, n$:

Propose:

$$C' \leftarrow \{h \cup \{H_i : h_i\} : h \in C, h_i \sim p(h_i \mid h_{i-1})\}$$

Weight:

Compute weights $w(h) = p(e_i \mid h_i)$ for $h \in C'$

Resample:


$$C \leftarrow K \text{ particles drawn independently from } \frac{w(h)}{\sum_{h' \in C'} w(h')}$$

[demo: particleFiltering({K:100})]

Particle filtering: implementation

For filtering questions, can optimize:

- Keep only value of last H_i for each particle
- Store count for each unique particle





Particle filtering demo

[see web version]



Summary



$$\mathbb{P}(H_3 \mid E_1 = 0, E_2 = 2, E_3 = 2)$$

- Use particles to represent an approximate distribution

Propose (transitions)

Weight (emissions)

Resample

- Can scale to large number of locations (unlike forward-backward)
- Maintains better particle diversity (compared to beam search)



Overall Summary: Bayesian Networks II

- Probabilistic programs as equivalent to Bayesian Networks
- Gibbs sampling is an algorithm for estimating marginal probabilities
- Forward Backward algorithm: Dynamic programming for inference (filtering and smoothing)
- Particle Filtering: Approximate inference for HMMs with large domains
- Next: learning the parameters of Bayesian networks